

# Time Series and Sequence Learning

## Lecture 8 – Parameter Estimation in Non-Linear/Non-Gaussian State Space Models

---

Johan Alenlöv, Linköping University

2020-09-28

## Summary of Lecture 7

---

# General State-Space Models

**Def.** A **General State-Space** model is given by:

$$\alpha_t \mid \alpha_{t-1} \sim q(\alpha_t \mid \alpha_{t-1})$$

$$y_t \mid \alpha_t \sim g(y_t \mid \alpha_t)$$

and initial distribution  $\alpha_1 \sim q(\alpha_1)$ .

**Thm.** The **joint-smoothing distribution** is given by

$$p(\alpha_{1:n} \mid y_{1:n}) = \frac{q(\alpha_1)g(y_1 \mid \alpha_1) \prod_{i=2}^n q(\alpha_i \mid \alpha_{i-1})g(y_i \mid \alpha_i)}{L_n(y_{1:n})},$$

where  $L_n(y_{1:n}) = \int q(\alpha_1)g(y_1 \mid \alpha_1) \prod_{i=2}^n q(\alpha_i \mid \alpha_{i-1})g(y_i \mid \alpha_i) d\alpha_{1:n}$  is the **likelihood**.

# Sequential Importance Sampling

- A first idea was to use **importance sampling** targeting the **joint-smoothing distribution**.
- We go from time  $n$  to  $n + 1$  in the following way:
  - Draw  $\alpha_{n+1}^i \sim f(\alpha_{n+1} | \alpha_{1:n}^i)$  (propagating)
  - Set  $\alpha_{1:n+1}^i = (\alpha_{1:n}^i, \alpha_{n+1}^i)$
  - Set  $\omega_{n+1}^i = \frac{q(\alpha_{n+1}^i | \alpha_n^i) g(y_{n+1} | \alpha_{n+1}^i)}{f(\alpha_{n+1}^i | \alpha_{1:n}^i)} \times \omega_n^i$ . (weighting)
- Simplest choice:  $f(\alpha_{n+1} | \alpha_{1:n}^i) = q(\alpha_{n+1} | \alpha_n^i)$
- We estimate using:

$$\sum_{i=1}^N \frac{\omega_{n+1}^i}{\Omega_{n+1}} h(\alpha_{n+1}^i) \approx \mathbb{E}[h(\alpha_{n+1}) | y_{1:n+1}]$$


$$N^{-1} \Omega_{n+1} = \frac{1}{N} \sum_{i=1}^N \omega_{n+1}^i \approx L(y_{1:n+1})$$

## Example: Linear Gaussian State Space Model

**Def.** A **Linear Gaussian State-Space (LGSS)** model is given by:

$$\begin{aligned}\alpha_t &= T\alpha_{t-1} + R\eta_t, & \eta_t &\sim \mathcal{N}(0, Q), \\ y_t &= Z\alpha_t + \varepsilon_t & \varepsilon_t &\sim \mathcal{N}(0, \sigma_\varepsilon^2),\end{aligned}$$

and initial distribution  $\alpha_1 \sim \mathcal{N}(a_1, P_1)$ .




# Example: Linear Gaussian State Space Model

**Def.** A **Linear Gaussian State-Space (LGSS)** model is given by:

$$\alpha_t = T\alpha_{t-1} + R\eta_t, \quad \eta_t \sim \mathcal{N}(0, Q),$$

$$y_t = Z\alpha_t + \varepsilon_t \quad \varepsilon_t \sim \mathcal{N}(0, \sigma_\varepsilon^2),$$

and initial distribution  $\alpha_1 \sim \mathcal{N}(a_1, P_1)$ .



# Sequential Importance Sampling with Resampling

- We solve the problem of **weight degeneracy** by **resampling** the particles!
- The most natural selection is to draw new particles  $(\tilde{\alpha}_{1:n}^i)_{i=1}^N$  among the **SIS** produced  $(\alpha_{1:n}^i)_{i=1}^N$  with probabilities given by the **normalized importance weights**.
- For  $i = 1, 2, \dots, N$  we let

$$\tilde{\alpha}_{1:n}^i = \alpha_{1:n}^j \quad \text{w. pr.} \quad \frac{\omega_n^j}{\Omega_n}$$

- Add this step **before** propagating the particles!


# Example: Linear Gaussian State Space Model

**Def.** A **Linear Gaussian State-Space (LGSS)** model is given by:

$$\alpha_t = T\alpha_{t-1} + R\eta_t, \quad \eta_t \sim \mathcal{N}(0, Q),$$

$$y_t = Z\alpha_t + \varepsilon_t \quad \varepsilon_t \sim \mathcal{N}(0, \sigma_\varepsilon^2),$$

and initial distribution  $\alpha_1 \sim \mathcal{N}(a_1, P_1)$ .





# Algorithm: Particle Filter

---

## Particle Filter:

---

Draw  $\alpha_1^i \sim f(\alpha_1)$

Set  $\omega_1^i = \frac{q(\alpha_1^i)g(y_1 | \alpha_1^i)}{f(\alpha_1^i)}$

Set  $\Omega_1 = \sum_{i=1}^N \omega_1^i$

**for**  $t = 2, 3, \dots, n$  **do**

    Draw  $l^j = j$  w. pr.  $\frac{\omega_{t-1}^j}{\Omega_{t-1}}$

    Draw  $\alpha_t^i \sim f(\alpha_t | \alpha_{1:t-1}^{l^j})$

    Set  $\omega_t^i = \frac{q(\alpha_t^i | \alpha_{t-1}^{l^j})g(y_t | \alpha_t^i)}{f(\alpha_t^i | \alpha_{1:t-1}^{l^j})}$

    Set  $\alpha_{1:t}^i = (\alpha_{1:t-1}^{l^j}, \alpha_t^i)$

    Set  $\Omega_t = \sum_{i=1}^N \omega_t^i$

**end for**

---

# Algorithm: Bootstrap Particle Filter

---

## Bootstrap Particle Filter:

---

Draw  $\alpha_1^i \sim q(\alpha_1)$

Set  $\omega_1^i = g(y_1 | \alpha_1^i)$

Set  $\Omega_1 = \sum_{i=1}^N \omega_1^i$

**for**  $t = 2, 3, \dots, n$  **do**

    Draw  $l^i = j$  w. pr.  $\frac{\omega_{t-1}^j}{\Omega_{t-1}}$

    Draw  $\alpha_t^i \sim q(\alpha_t | \alpha_{t-1}^{l^i})$

    Set  $\omega_t^i = g(y_t | \alpha_t^i)$

    Set  $\alpha_{1:t}^i = (\alpha_{1:t-1}^{l^i}, \alpha_t^i)$

    Set  $\Omega_t = \sum_{i=1}^N \omega_t^i$

**end for**

---

# Sequential Importance Sampling with Resampling


In Python:

---


```
1 part = np.zeros((n,N))
2 logwgt = np.zeros((n,N))
3 wgt = np.zeros((n,N))
4 ests = np.zeros(n)
5 part[0,:] = np.random.randn(N)
6 logwgt[0,:] = logwgtfun(xpart[0,:],y[0])
7 wgt[0,:] = np.exp(logwgt[0,:])
8 ests[0] = np.sum(wgt[0,:] * part[0,:])/np.sum(wgt[i+1,:])
9 for i in range(n-1):
10     ind = np.random.choice(N, size=N, replace=True, p=wgt[i,:])
11     part[i+1,:] = a*part[i,ind] + np.random.randn(N)
12     logwgt[i+1,:] = logwgtfun(xpart[i+1,:],y[i+1])
13     wgt[i+1,:] = np.exp(logwgt[i+1,:])
14     ests[i+1] = np.sum(wgt[i+1,:] * part[i+1,:])/np.sum(wgt[i+1,:])
```

---


# Particle Filter Movie




# Particle Filter Movie




# Particle Filter Movie



# Particle Filter Movie




# Particle Filter Movie







# Particle Filter Movie




# Particle Filter Movie




# Particle Filter Movie




# Particle Filter Movie




# Particle Filter Movie




# Particle Filter Movie



# Particle Filter Movie




# Particle Filter Movie







# Particle Filter Movie




# Particle Filter Movie




# Particle Filter Movie




# Particle Filter Movie




# Particle Filter Movie




# Particle Filter Movie



# Particle Filter Movie




# Particle Filter Movie







# Particle Filter Movie



# Particle Filter Movie



# Particle Filter Movie



# Parameter Estimation in General State-Space Models

---

# Calculating the Log-Likelihood

- For the **Sequential Importance Sampling** algorithm we had that

$$\frac{1}{N} \sum_{i=1}^N \omega_n^i \approx L(y_{1:n})$$

- When adding resampling this is **not** a valid estimator for the likelihood.
- Instead we use

$$\prod_{t=1}^n \left( \frac{1}{N} \sum_{i=1}^N \omega_t^i \right) \approx L(y_{1:n})$$


- **Note:**  $\omega_t^i$  are the **unnormalized** weights!
- This is an **unbiased** estimator of the likelihood!

## Example: Gaussian State Space Model


**Linear Gaussian State Space Model** We look at the model defined by the following equations:

$$\begin{cases} \alpha_{t+1} = a\alpha_t + \eta_{t+1}, & \eta_{t+1} \sim \mathcal{N}(0, 1), \\ y_t = \alpha_t + \varepsilon_t, & \varepsilon_t \sim \mathcal{N}(0, 1), \end{cases}$$


with initial distribution  $\alpha_1 \sim \mathcal{N}(0, 1/(1 - a^2))$



## Example: Gaussian State Space Model

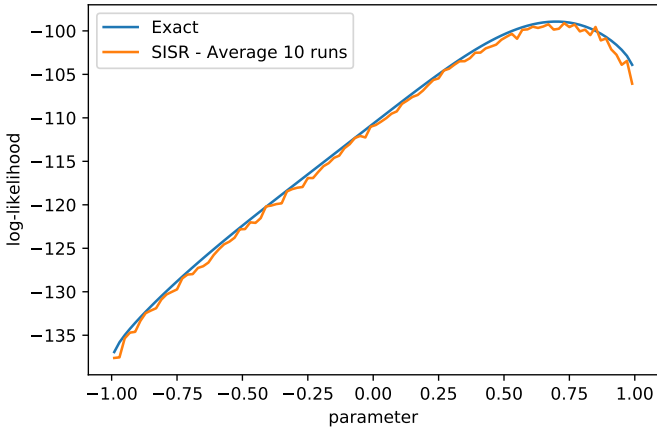


# Example: Gaussian State Space Model





# Example: Gaussian State Space Model



# Expectation-Maximization Algorithm

- As in the LGSS model we can use **expectation-maximization** algorithm for parameter estimation.
- The **E-step** is calculating:

$$\begin{aligned}\mathcal{Q}(\tilde{\theta}, \theta) &= \mathbb{E}[\log p_{\theta}(\alpha_{1:n}, y_{1:n}) \mid y_{1:n}, \tilde{\theta}] \\ &= \mathbb{E}[\log q_{\theta}(\alpha_1) \mid y_{1:n}, \tilde{\theta}] + \sum_{t=2}^n \mathbb{E}[\log q_{\theta}(\alpha_t \mid \alpha_{t-1}) \mid y_{1:n}, \tilde{\theta}] \\ &\quad + \sum_{t=1}^n \mathbb{E}[\log g_{\theta}(y_t \mid \alpha_t) \mid y_{1:n}, \tilde{\theta}]\end{aligned}$$

- The **M-step** is to maximize  $\mathcal{Q}(\tilde{\theta}, \theta)$ :

$$\theta^* = \arg \max_{\theta} \mathcal{Q}(\tilde{\theta}, \theta)$$

## Interlude I: Exponential Family

**Def:** A distribution belongs to the **exponential family** if its density function can be written as,

$$p_{\theta}(x) = h(x) \exp(\mathbf{n}(\theta) \cdot \mathbf{T}(x) - A(\theta)),$$

where

- $\mathbf{T}(x)$  is a **sufficient statistic**
- $\mathbf{n}(\theta)$  is the **natural parameter**

**Ex:** A **Gaussian** distribution with mean  $\mu$  and variance  $\sigma^2$  belongs to the **exponential family** using

$$\underbrace{\frac{1}{\sqrt{2\pi}}}_{h(x)} \exp \left( \underbrace{\begin{pmatrix} \frac{\mu}{\sigma^2} \\ -\frac{1}{2\sigma^2} \end{pmatrix}}_{\mathbf{n}(\theta)} \cdot \underbrace{\begin{pmatrix} x \\ x^2 \end{pmatrix}}_{\mathbf{T}(x)} - \underbrace{\left( \frac{\mu^2}{2\sigma^2} + \frac{1}{2} \log(\sigma^2) \right)}_{A(\theta)} \right)$$

# Expectation-Maximization Algorithm II

- We assume that  $q_{\theta}$  and  $g_{\theta}$  belongs to the **exponential family**.
- The **E-step** now reduces to,

## Expectation-Maximization Algorithm II

- We assume that  $q_{\theta}$  and  $g_{\theta}$  belongs to the **exponential family**.
- The **E-step** now reduces to,

$$\begin{aligned} Q(\tilde{\theta}, \theta) &= \mathbb{E}[\log h_q^1(\alpha_1) + \mathbf{n}_q^1(\theta) \cdot \mathbf{T}_q(\alpha_1) - A_q^1(\theta) \mid y_{1:n}, \tilde{\theta}] \\ &\quad + \sum_{t=2}^n \mathbb{E}[\log h_q(\alpha_t, \alpha_{t-1}) + \mathbf{n}_q(\theta) \cdot \mathbf{T}_q(\alpha_t, \alpha_{t-1}) - A_q(\theta) \mid y_{1:n}, \tilde{\theta}] \\ &\quad + \sum_{t=1}^n \mathbb{E}[\log h_g(y_t, \alpha_t) + \mathbf{n}_g(\theta) \cdot \mathbf{T}_g(y_t, \alpha_t) - A_g(\theta) \mid y_{1:n}, \tilde{\theta}] \end{aligned}$$

## Expectation-Maximization Algorithm II

- We assume that  $q_{\theta}$  and  $g_{\theta}$  belongs to the **exponential family**.
- The **E-step** now reduces to,

$$\begin{aligned} Q(\tilde{\theta}, \theta) &= \mathbb{E}[\log h_q^1(\alpha_1) \mid y_{1:n}, \tilde{\theta}] + \mathbf{n}_q^1(\theta) \cdot \mathbb{E}[\mathbf{T}_q(\alpha_1) \mid y_{1:n}, \tilde{\theta}] - A_q^1(\theta) \\ &\quad + \sum_{t=2}^n \mathbb{E}[\log h_q(\alpha_t, \alpha_{t-1}) \mid y_{1:n}, \tilde{\theta}] + \mathbf{n}_q(\theta) \cdot \mathbb{E}[\mathbf{T}_q(\alpha_t, \alpha_{t-1}) \mid y_{1:n}, \tilde{\theta}] \\ &\quad \quad \quad - A_q(\theta) \\ &\quad + \sum_{t=1}^n \mathbb{E}[\log h_g(y_t, \alpha_t) \mid y_{1:n}, \tilde{\theta}] + \mathbf{n}_g(\theta) \cdot \mathbb{E}[\mathbf{T}_g(y_t, \alpha_t) \mid y_{1:n}, \tilde{\theta}] \\ &\quad \quad \quad - A_g(\theta) \end{aligned}$$

## Expectation-Maximization Algorithm II

- We assume that  $q_{\theta}$  and  $g_{\theta}$  belongs to the **exponential family**.
- The **E-step** now reduces to,  
Calculate the **smoothed sufficient statistics**,

$$\mathbf{T}_1 = \mathbb{E}[\mathbf{T}_{\mathbf{q}}(\alpha_1) \mid y_{1:n}, \tilde{\theta}],$$

$$\mathbf{T}_2 = \sum_{t=2}^n \mathbb{E}[\mathbf{T}_{\mathbf{q}}(\alpha_t, \alpha_{t-1}) \mid y_{1:n}, \tilde{\theta}],$$

$$\mathbf{T}_3 = \sum_{t=1}^n \mathbb{E}[\mathbf{T}_{\mathbf{g}}(y_t, \alpha_t) \mid y_{1:n}, \tilde{\theta}].$$

## Expectation-Maximization Algorithm II

- We assume that  $q_{\theta}$  and  $g_{\theta}$  belongs to the **exponential family**.
- The **E-step** now reduces to,  
Calculate the **smoothed sufficient statistics**,

$$\mathbf{T}_1 = \mathbb{E}[\mathbf{T}_q(\alpha_1) \mid y_{1:n}, \tilde{\theta}],$$

$$\mathbf{T}_2 = \sum_{t=2}^n \mathbb{E}[\mathbf{T}_q(\alpha_t, \alpha_{t-1}) \mid y_{1:n}, \tilde{\theta}],$$

$$\mathbf{T}_3 = \sum_{t=1}^n \mathbb{E}[\mathbf{T}_g(y_t, \alpha_t) \mid y_{1:n}, \tilde{\theta}].$$


- The **M-step** becomes maximize,

$$\mathbf{n}_q^1(\theta) \cdot \mathbf{T}_1 - A_q^1(\theta) + \mathbf{n}_q(\theta) \cdot \mathbf{T}_2 - A_q(\theta) + \mathbf{n}_g(\theta) \cdot \mathbf{T}_3 - A_g(\theta)$$



# Joint-Smoothing Distribution

- As previously for the EM-algorithm we need the **smoothing distribution**.
- Good news! When we derived the particle filter we did it for the **joint-smoothing distribution**.
- The **particle trajectories** approximate the joint-smoothing distribution.
- **Idea:** Run the the particle filter, store the trajectories and approximate the smoothing distribution.
- **Problem:** Resampling collapses the trajectories.



# Fixed-Lag Smoothing

- There are many algorithms to perform **smoothing** in general State Space models.
- Most of them involve intricate backward passes or Markov Chains to estimate the smoothing distribution.
- A **simple** smoothing algorithm is the **fixed-lag smoothing**.
- We approximate using

$$\mathbb{E}[h(\alpha_t) \mid y_{1:n}] \approx \mathbb{E}[h(\alpha_t) \mid y_{1:t+\ell}],$$

where  $\ell$  is a **fixed lag**.

## Example: Stochastic Volatility

**Stochastic volatility model.** The model is defined by the equations,

$$\begin{cases} \alpha_k = a\alpha_{k-1} + \sigma_\eta \eta_k, & \eta_k \sim \mathcal{N}(0, s^2) \\ y_k = b \exp(\alpha_k/2) \varepsilon_k, & \varepsilon_k \sim \mathcal{N}(0, 1) \end{cases}$$

For simplicity we assume that  $\alpha_1 \sim \mathcal{N}(0, 1)$ .


**EM-algorithm**, in this model we find four **sufficient statistics**:

$$\begin{aligned} t_1 &= \sum_{t=2}^n \mathbb{E}[\alpha_t^2 \mid y_{1:n}] & t_2 &= \sum_{t=2}^n \mathbb{E}[\alpha_t \alpha_{t-1} \mid y_{1:n}] \\ t_3 &= \sum_{t=2}^n \mathbb{E}[\alpha_{t-1}^2 \mid y_{1:n}] & t_4 &= \sum_{t=1}^n \mathbb{E}[y_t^2 \exp(-\alpha_t) \mid y_{1:n}] \end{aligned}$$


The parameters are updated as,

$$a = \frac{t_2}{t_3} \quad s^2 = \frac{1}{n-1} \left( t_1 - \frac{t_2^2}{t_3} \right) \quad b^2 = \frac{t_4}{n}$$

## Example: Stochastic Volatility



## Example: Stochastic Volatility



# Extensions to Particle Filters

---

# Adaptive Resampling

- So far we have done **resampling every iteration**.
- This is **unnecessary**, we only need to do it if the weights have degenerated.
- Use **effective sample size** (ESS) as a measure for degeneracy,

$$\text{ESS}_t = \frac{(\sum_{i=1}^N \omega_t^i)^2}{\sum_{i=1}^N (\omega_t^i)^2}$$

- If all weights are **equal** then  $\text{ESS}_t = N$ .
- If all weights except one is zero then  $\text{ESS}_t = 1$ .
- Choose a **threshold**  $N_{\text{ESS}}$ , if  $\text{ESS}_t < N_{\text{ESS}}$  then resample, otherwise no resampling happens.
- If no resampling happens the weights should be updated as in the SIS algorithm.

# Algorithm: Particle Filter with Adaptive Resampling

---

## Particle Filter:

---

Draw  $\alpha_1^i \sim f(\alpha_1)$

Set  $\omega_1^i = \frac{q(\alpha_1^i)g(y_1 | \alpha_1^i)}{f(\alpha_1^i)}$

Set  $\Omega_1 = \sum_{i=1}^N \omega_1^i$

**for**  $t = 2, 3, \dots, n$  **do**

Calculate  $\text{ESS}_{t-1} = \frac{(\sum_{i=1}^N \omega_{t-1}^i)^2}{\sum_{i=1}^N (\omega_{t-1}^i)^2}$

**if**  $\text{ESS}_{t-1} < N_{\text{ESS}}$  **then**

Draw  $l' = j$  w. pr.  $\frac{\omega_{t-1}^j}{\Omega_{t-1}}$

Set  $\alpha_{1:t}^i = (\alpha_{1:t-1}^{l'}, \alpha_t^i)$

Set  $\omega_{t-1}^i = 1/N$

**end if**

Draw  $\alpha_t^i \sim f(\alpha_t | \alpha_{1:t-1}^i)$

Set  $\omega_t^i = \frac{q(\alpha_t^i | \alpha_{t-1}^i)g(y_t | \alpha_t^i)}{f(\alpha_t^i | \alpha_{1:t-1}^i)} \times \omega_{t-1}^i$


Set  $\Omega_t = \sum_{i=1}^N \omega_t^i$

**end for**

---




## Example: Adaptive resampling



Resampling: 43%

## Example: Adaptive resampling



Resampling: 36%

# Different Resampling Schemes

- We have so far only considered **multinomial resampling**.
- There are many different schemes,
  - **Residual resampling**
  - **Stratified resampling**
  - **Systematic resampling**

# Summary

- We looked at how to calculate the **likelihood** using a **particle filter**.
  - The randomness of the algorithm gives us **noisy** likelihood estimate.
  - **Solve** by performing many runs and average.
- We looked at the **EM-algorithm**.
  - Easier if the model belongs to the **exponential family**.
  - Requires, again, the **smoothing distribution**.
  - Looked at the **fixed-lag smoother**.
- We looked at some extensions regarding the **resampling**.
  - By calculating the ESS we performed **adaptive resampling** by only resampling when needed.
  - We looked at some different resampling schemes.