

Time Series and Sequence Learning

Lecture 8 – Parameter Estimation in Non-Linear/Non-Gaussian State Space Models

Johan Alenlöv, Linköping University

2024-09-30

Summary of Lecture 7

General State-Space Models

Def. A **General State-Space** model is given by:

$$\alpha_t \mid \alpha_{t-1} \sim q(\alpha_t \mid \alpha_{t-1})$$

$$y_t \mid \alpha_t \sim g(y_t \mid \alpha_t)$$

and initial distribution $\alpha_1 \sim q(\alpha_1)$.

General State-Space Models

Def. A **General State-Space** model is given by:

$$\alpha_t \mid \alpha_{t-1} \sim q(\alpha_t \mid \alpha_{t-1})$$

$$y_t \mid \alpha_t \sim g(y_t \mid \alpha_t)$$

and initial distribution $\alpha_1 \sim q(\alpha_1)$.

Thm. The **joint-smoothing distribution** is given by

$$p(\alpha_{1:n} \mid y_{1:n}) = \frac{q(\alpha_1)g(y_1 \mid \alpha_1) \prod_{i=2}^n q(\alpha_i \mid \alpha_{i-1})g(y_i \mid \alpha_i)}{L_n(y_{1:n})},$$

where $L_n(y_{1:n}) = \int q(\alpha_1)g(y_1 \mid \alpha_1) \prod_{i=2}^n q(\alpha_i \mid \alpha_{i-1})g(y_i \mid \alpha_i) d\alpha_{1:n}$ is the **likelihood**.

Sequential Importance Sampling

- A first idea was to use **importance sampling** targeting the **joint-smoothing distribution**.

Sequential Importance Sampling

- A first idea was to use **importance sampling** targeting the **joint-smoothing distribution**.
- We go from time n to $n + 1$ in the following way:
 - Draw $\alpha_{n+1}^i \sim f(\alpha_{n+1} | \alpha_{1:n}^i)$ (propagating)
 - Set $\alpha_{1:n+1}^i = (\alpha_{1:n}^i, \alpha_{n+1}^i)$
 - Set $\omega_{n+1}^i = \frac{q(\alpha_{n+1}^i | \alpha_n^i) g(y_{n+1} | \alpha_{n+1}^i)}{f(\alpha_{n+1}^i | \alpha_{1:n}^i)} \times \omega_n^i$. (weighting)

Sequential Importance Sampling

- A first idea was to use **importance sampling** targeting the **joint-smoothing distribution**.
- We go from time n to $n + 1$ in the following way:
 - Draw $\alpha_{n+1}^i \sim f(\alpha_{n+1} | \alpha_{1:n}^i)$ (propagating)
 - Set $\alpha_{1:n+1}^i = (\alpha_{1:n}^i, \alpha_{n+1}^i)$
 - Set $\omega_{n+1}^i = \frac{q(\alpha_{n+1}^i | \alpha_n^i) g(y_{n+1} | \alpha_{n+1}^i)}{f(\alpha_{n+1}^i | \alpha_{1:n}^i)} \times \omega_n^i$. (weighting)
- Simplest choice: $f(\alpha_{n+1} | \alpha_{1:n}^i) = q(\alpha_{n+1} | \alpha_n^i)$

Sequential Importance Sampling

- A first idea was to use **importance sampling** targeting the **joint-smoothing distribution**.
- We go from time n to $n + 1$ in the following way:
 - Draw $\alpha_{n+1}^i \sim f(\alpha_{n+1} | \alpha_{1:n}^i)$ (propagating)
 - Set $\alpha_{1:n+1}^i = (\alpha_{1:n}^i, \alpha_{n+1}^i)$
 - Set $\omega_{n+1}^i = \frac{q(\alpha_{n+1}^i | \alpha_n^i) g(y_{n+1} | \alpha_{n+1}^i)}{f(\alpha_{n+1}^i | \alpha_{1:n}^i)} \times \omega_n^i$. (weighting)
- Simplest choice: $f(\alpha_{n+1} | \alpha_{1:n}^i) = q(\alpha_{n+1} | \alpha_n^i)$
- We estimate using:

$$\sum_{i=1}^N \frac{\omega_{n+1}^i}{\Omega_{n+1}} h(\alpha_{n+1}^i) \approx \mathbb{E}[h(\alpha_{n+1}) | y_{1:n+1}]$$


$$N^{-1} \Omega_{n+1} = \frac{1}{N} \sum_{i=1}^N \omega_{n+1}^i \approx L(y_{1:n+1})$$

Example: Linear Gaussian State Space Model

Def. A **Linear Gaussian State-Space (LGSS)** model is given by:

$$\begin{aligned}\alpha_t &= T\alpha_{t-1} + R\eta_t, & \eta_t &\sim \mathcal{N}(0, Q), \\ y_t &= Z\alpha_t + \varepsilon_t & \varepsilon_t &\sim \mathcal{N}(0, \sigma_\varepsilon^2),\end{aligned}$$

and initial distribution $\alpha_1 \sim \mathcal{N}(a_1, P_1)$.




Example: Linear Gaussian State Space Model

Def. A **Linear Gaussian State-Space (LGSS)** model is given by:

$$\begin{aligned}\alpha_t &= T\alpha_{t-1} + R\eta_t, & \eta_t &\sim \mathcal{N}(0, Q), \\ y_t &= Z\alpha_t + \varepsilon_t & \varepsilon_t &\sim \mathcal{N}(0, \sigma_\varepsilon^2),\end{aligned}$$

and initial distribution $\alpha_1 \sim \mathcal{N}(a_1, P_1)$.



Sequential Importance Sampling with Resampling

- We solve the problem of **weight degeneracy** by **resampling** the particles!

Sequential Importance Sampling with Resampling

- We solve the problem of **weight degeneracy** by **resampling** the particles!
- The most natural selection is to draw new particles $(\tilde{\alpha}_{1:n}^i)_{i=1}^N$ among the **SIS** produced $(\alpha_{1:n}^i)_{i=1}^N$ with probabilities given by the **normalized importance weights**.

Sequential Importance Sampling with Resampling

- We solve the problem of **weight degeneracy** by **resampling** the particles!
- The most natural selection is to draw new particles $(\tilde{\alpha}_{1:n}^i)_{i=1}^N$ among the **SIS** produced $(\alpha_{1:n}^i)_{i=1}^N$ with probabilities given by the **normalized importance weights**.
- For $i = 1, 2, \dots, N$ we let

$$\tilde{\alpha}_{1:n}^i = \alpha_{1:n}^j \quad \text{w. pr.} \quad \frac{\omega_n^j}{\Omega_n}$$

Sequential Importance Sampling with Resampling

- We solve the problem of **weight degeneracy** by **resampling** the particles!
- The most natural selection is to draw new particles $(\tilde{\alpha}_{1:n}^i)_{i=1}^N$ among the **SIS** produced $(\alpha_{1:n}^i)_{i=1}^N$ with probabilities given by the **normalized importance weights**.
- For $i = 1, 2, \dots, N$ we let

$$\tilde{\alpha}_{1:n}^i = \alpha_{1:n}^j \quad \text{w. pr.} \quad \frac{\omega_n^j}{\Omega_n}$$


- Add this step **before** propagating the particles!

Example: Linear Gaussian State Space Model

Def. A **Linear Gaussian State-Space (LGSS)** model is given by:

$$\begin{aligned}\alpha_t &= T\alpha_{t-1} + R\eta_t, & \eta_t &\sim \mathcal{N}(0, Q), \\ y_t &= Z\alpha_t + \varepsilon_t & \varepsilon_t &\sim \mathcal{N}(0, \sigma_\varepsilon^2),\end{aligned}$$

and initial distribution $\alpha_1 \sim \mathcal{N}(a_1, P_1)$.



Algorithm: Particle Filter

Particle Filter:

Draw $\alpha_1^i \sim f(\alpha_1)$

Set $\omega_1^i = \frac{q(\alpha_1^i)g(y_1 | \alpha_1^i)}{f(\alpha_1^i)}$

Set $\Omega_1 = \sum_{i=1}^N \omega_1^i$

for $t = 2, 3, \dots, n$ **do**

 Draw $l^i = j$ w. pr. $\frac{\omega_{t-1}^j}{\Omega_{t-1}}$

 Draw $\alpha_t^i \sim f(\alpha_t | \alpha_{1:t-1}^{l^i})$

 Set $\omega_t^i = \frac{q(\alpha_t^i | \alpha_{t-1}^{l^i})g(y_t | \alpha_t^i)}{f(\alpha_t^i | \alpha_{1:t-1}^{l^i})}$

 Set $\alpha_{1:t}^i = (\alpha_{1:t-1}^{l^i}, \alpha_t^i)$

 Set $\Omega_t = \sum_{i=1}^N \omega_t^i$

end for

Algorithm: Bootstrap Particle Filter

Bootstrap Particle Filter:

Draw $\alpha_1^i \sim q(\alpha_1)$

Set $\omega_1^i = g(y_1 | \alpha_1^i)$

Set $\Omega_1 = \sum_{i=1}^N \omega_1^i$

for $t = 2, 3, \dots, n$ **do**

 Draw $l^i = j$ w. pr. $\frac{\omega_{t-1}^j}{\Omega_{t-1}}$

 Draw $\alpha_t^i \sim q(\alpha_t | \alpha_{t-1}^{l^i})$

 Set $\omega_t^i = g(y_t | \alpha_t^i)$

 Set $\alpha_{1:t}^i = (\alpha_{1:t-1}^{l^i}, \alpha_t^i)$

 Set $\Omega_t = \sum_{i=1}^N \omega_t^i$

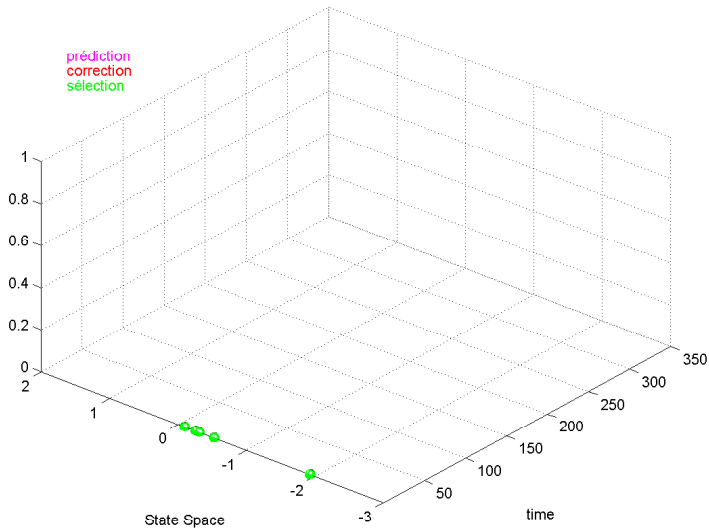
end for

Sequential Importance Sampling with Resampling


In Python:

```
1 part = np.zeros((n,N))
2 logwgt = np.zeros((n,N))
3 wgt = np.zeros((n,N))
4 ests = np.zeros(n)
5 part[0,:] = np.random.randn(N)
6 logwgt[0,:] = logwgtfun(xpart[0,:],y[0])
7 wgt[0,:] = np.exp(logwgt[0,:])
8 ests[0] = np.sum(wgt[0,:] * part[0,:])/np.sum(wgt[i+1,:])
9 for i in range(n-1):
10     ind = np.random.choice(N, size=N, replace=True, p=wgt[i,:])
11     part[i+1,:] = a*part[i,ind] + np.random.randn(N)
12     logwgt[i+1,:] = logwgtfun(xpart[i+1,:],y[i+1])
13     wgt[i+1,:] = np.exp(logwgt[i+1,:])
14     ests[i+1] = np.sum(wgt[i+1,:] * part[i+1,:])/np.sum(wgt[i+1,:])
```


Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie




Particle Filter Movie



Particle Filter Movie



Particle Filter Movie



Parameter Estimation in General State-Space Models

Calculating the Log-Likelihood

- For the **Sequential Importance Sampling** algorithm we have that

$$\frac{1}{N} \sum_{i=1}^N \omega_n^i \approx L(y_{1:n})$$

- When adding resampling this is **not** a valid estimator for the likelihood.

Calculating the Log-Likelihood

- For the **Sequential Importance Sampling** algorithm we have that

$$\frac{1}{N} \sum_{i=1}^N \omega_n^i \approx L(y_{1:n})$$

- When adding resampling this is **not** a valid estimator for the likelihood.
- Instead we use

$$\prod_{t=1}^n \left(\frac{1}{N} \sum_{i=1}^N \omega_t^i \right) \approx L(y_{1:n})$$

- Note:** ω_t^i are the **unnormalized** weights!

Calculating the Log-Likelihood

- For the **Sequential Importance Sampling** algorithm we have that

$$\frac{1}{N} \sum_{i=1}^N \omega_n^i \approx L(y_{1:n})$$

- When adding resampling this is **not** a valid estimator for the likelihood.
- Instead we use

$$\prod_{t=1}^n \left(\frac{1}{N} \sum_{i=1}^N \omega_t^i \right) \approx L(y_{1:n})$$


- Note:** ω_t^i are the **unnormalized** weights!
- The is an **unbiased** estimator of the likelihood!

Example: Gaussian State Space Model


Linear Gaussian State Space Model We look at the model defined by the following equations:

$$\begin{cases} \alpha_{t+1} = a\alpha_t + \eta_{t+1}, & \eta_{t+1} \sim \mathcal{N}(0, 1), \\ y_t = \alpha_t + \varepsilon_t, & \varepsilon_t \sim \mathcal{N}(0, 1), \end{cases}$$


with initial distribution $\alpha_1 \sim \mathcal{N}(0, 1/(1 - a^2))$




Example: Gaussian State Space Model



Example: Gaussian State Space Model



Example: Gaussian State Space Model



Expectation-Maximization

Expectation-Maximization Algorithm

- As in the LGSS model we can use **expectation-maximization** algorithm for parameter estimation.

Expectation-Maximization Algorithm

- As in the LGSS model we can use **expectation-maximization** algorithm for parameter estimation.
- The **E-step** is calculating:

$$\begin{aligned} Q(\tilde{\theta}, \theta) &= \mathbb{E}[\log p_{\theta}(\alpha_{1:n}, y_{1:n}) \mid y_{1:n}, \tilde{\theta}] \\ &= \mathbb{E}[\log q_{\theta}(\alpha_1) \mid y_{1:n}, \tilde{\theta}] + \sum_{t=2}^n \mathbb{E}[\log q_{\theta}(\alpha_t \mid \alpha_{t-1}) \mid y_{1:n}, \tilde{\theta}] \\ &\quad + \sum_{t=1}^n \mathbb{E}[\log g_{\theta}(y_t \mid \alpha_t) \mid y_{1:n}, \tilde{\theta}] \end{aligned}$$

Expectation-Maximization Algorithm

- As in the LGSS model we can use **expectation-maximization** algorithm for parameter estimation.
- The **E-step** is calculating:

$$\begin{aligned} Q(\tilde{\theta}, \theta) &= \mathbb{E}[\log p_{\theta}(\alpha_{1:n}, y_{1:n}) \mid y_{1:n}, \tilde{\theta}] \\ &= \mathbb{E}[\log q_{\theta}(\alpha_1) \mid y_{1:n}, \tilde{\theta}] + \sum_{t=2}^n \mathbb{E}[\log q_{\theta}(\alpha_t \mid \alpha_{t-1}) \mid y_{1:n}, \tilde{\theta}] \\ &\quad + \sum_{t=1}^n \mathbb{E}[\log g_{\theta}(y_t \mid \alpha_t) \mid y_{1:n}, \tilde{\theta}] \end{aligned}$$

- The **M-step** is to maximize $Q(\tilde{\theta}, \theta)$:

$$\theta^* = \arg \max_{\theta} Q(\tilde{\theta}, \theta)$$

Interlude I: Exponential Family

Def: A distribution belongs to the **exponential family** if its density function can be written as,

$$p_{\theta}(x) = h(x) \exp (n(\theta) \cdot T(x) - A(\theta)),$$

where

- $T(x)$ is a **sufficient statistic**
- $n(\theta)$ is the **natural parameter**

Interlude I: Exponential Family

Def: A distribution belongs to the **exponential family** if its density function can be written as,

$$p_{\theta}(x) = h(x) \exp (n(\theta) \cdot T(x) - A(\theta)),$$

where

- $T(x)$ is a **sufficient statistic**
- $n(\theta)$ is the **natural parameter**

Ex: A **Gaussian** distribution with mean μ and variance σ^2 belongs to the **exponential family** using

$$\underbrace{\frac{1}{\sqrt{2\pi}}}_{h(x)} \exp \left(\underbrace{\begin{pmatrix} \frac{\mu}{\sigma^2} \\ -\frac{1}{2\sigma^2} \end{pmatrix}}_{n(\theta)} \cdot \underbrace{\begin{pmatrix} x \\ x^2 \end{pmatrix}}_{T(x)} - \underbrace{\left(\frac{\mu^2}{2\sigma^2} + \frac{1}{2} \log(\sigma^2) \right)}_{A(\theta)} \right)$$

- We assume that q_{θ} and g_{θ} belongs to the exponential family.

Expectation-Maximization Algorithm II

- We assume that q_{θ} and g_{θ} belongs to the **exponential family**.
- The **E-step** now reduces to,

$$\begin{aligned} Q(\tilde{\theta}, \theta) &= \mathbb{E}[\log h_q^1(\alpha_1) + \mathbf{n}_q^1(\theta) \cdot \mathbf{T}_q(\alpha_1) - A_q^1(\theta) \mid y_{1:n}, \tilde{\theta}] \\ &\quad + \sum_{t=2}^n \mathbb{E}[\log h_q(\alpha_t, \alpha_{t-1}) + \mathbf{n}_q(\theta) \cdot \mathbf{T}_q(\alpha_t, \alpha_{t-1}) - A_q(\theta) \mid y_{1:n}, \tilde{\theta}] \\ &\quad + \sum_{t=1}^n \mathbb{E}[\log h_g(y_t, \alpha_t) + \mathbf{n}_g(\theta) \cdot \mathbf{T}_g(y_t, \alpha_t) - A_g(\theta) \mid y_{1:n}, \tilde{\theta}] \end{aligned}$$

Expectation-Maximization Algorithm II

- We assume that q_{θ} and g_{θ} belongs to the **exponential family**.
- The **E-step** now reduces to,

$$\begin{aligned} \mathcal{Q}(\tilde{\theta}, \theta) &= \mathbb{E}[\log h_q^1(\alpha_1) | y_{1:n}, \tilde{\theta}] + \mathbf{n}_q^1(\theta) \cdot \mathbb{E}[\mathbf{T}_q(\alpha_1) | y_{1:n}, \tilde{\theta}] - A_q^1(\theta) \\ &\quad + \sum_{t=2}^n \mathbb{E}[\log h_q(\alpha_t, \alpha_{t-1}) | y_{1:n}, \tilde{\theta}] + \mathbf{n}_q(\theta) \cdot \mathbb{E}[\mathbf{T}_q(\alpha_t, \alpha_{t-1}) | y_{1:n}, \tilde{\theta}] \\ &\quad \quad \quad - A_q(\theta) \\ &\quad + \sum_{t=1}^n \mathbb{E}[\log h_g(y_t, \alpha_t) | y_{1:n}, \tilde{\theta}] + \mathbf{n}_g(\theta) \cdot \mathbb{E}[\mathbf{T}_g(y_t, \alpha_t) | y_{1:n}, \tilde{\theta}] \\ &\quad \quad \quad - A_g(\theta) \end{aligned}$$

Expectation-Maximization Algorithm II

- We assume that q_{θ} and g_{θ} belongs to the **exponential family**.
- The **E-step** now reduces to,
Calculate the **smoothed sufficient statistics**,

$$T_1 = \mathbb{E}[T_q(\alpha_1) \mid y_{1:n}, \tilde{\theta}],$$

$$T_2 = \sum_{t=2}^n \mathbb{E}[T_q(\alpha_t, \alpha_{t-1}) \mid y_{1:n}, \tilde{\theta}],$$

$$T_3 = \sum_{t=1}^n \mathbb{E}[T_g(y_t, \alpha_t) \mid y_{1:n}, \tilde{\theta}].$$

Expectation-Maximization Algorithm II

- We assume that q_{θ} and g_{θ} belongs to the **exponential family**.
- The **E-step** now reduces to,
Calculate the **smoothed sufficient statistics**,

$$T_1 = \mathbb{E}[T_q(\alpha_1) \mid y_{1:n}, \tilde{\theta}],$$

$$T_2 = \sum_{t=2}^n \mathbb{E}[T_q(\alpha_t, \alpha_{t-1}) \mid y_{1:n}, \tilde{\theta}],$$

$$T_3 = \sum_{t=1}^n \mathbb{E}[T_g(y_t, \alpha_t) \mid y_{1:n}, \tilde{\theta}].$$

- The **M-step** becomes maximize,

$$n_q^1(\theta) \cdot T_1 - A_q^1(\theta) + n_q(\theta) \cdot T_2 - A_q(\theta) + n_g(\theta) \cdot T_3 - A_g(\theta)$$

Smoothing

Joint-Smoothing Distribution

- As previously for the EM-algorithm we need the **smoothing distribution**.

Joint-Smoothing Distribution

- As previously for the EM-algorithm we need the **smoothing distribution**.
- Good news! When we derived the particle filter we did it for the **joint-smoothing distribution**.

Joint-Smoothing Distribution


- As previously for the EM-algorithm we need the **smoothing distribution**.
- Good news! When we derived the particle filter we did it for the **joint-smoothing distribution**.
- The **particle trajectories** approximate the joint-smoothing distribution.

Joint-Smoothing Distribution

- As previously for the EM-algorithm we need the **smoothing distribution**.
- Good news! When we derived the particle filter we did it for the **joint-smoothing distribution**.
- The **particle trajectories** approximate the joint-smoothing distribution.
- **Idea:** Run the the particle filter, store the trajectories and approximate the smoothing distribution.

Joint-Smoothing Distribution

- As previously for the EM-algorithm we need the **smoothing distribution**.
- Good news! When we derived the particle filter we did it for the **joint-smoothing distribution**.
- The **particle trajectories** approximate the joint-smoothing distribution.
- **Idea:** Run the the particle filter, store the trajectories and approximate the smoothing distribution.
- **Problem:** Resampling collapses the trajectories.



Fixed-Lag Smoothing

- There are many algorithms to perform **smoothing** in general State Space models.
- Most of them involve intricate backward passes or Markov Chains to estimate the smoothing distribution.

Fixed-Lag Smoothing

- There are many algorithms to perform **smoothing** in general State Space models.
- Most of them involve intricate backward passes or Markov Chains to estimate the smoothing distribution.
- A **simple** smoothing algorithm is the **fixed-lag smoothing**.
- We approximate using

$$\mathbb{E}[h(\alpha_t) | y_{1:n}] \approx \mathbb{E}[h(\alpha_t) | y_{1:t+\ell}],$$

where ℓ is a **fixed lag**.

Example: Stochastic Volatility

Stochastic volatility model. The model is defined by the equations,

$$\begin{cases} \alpha_k = a\alpha_{k-1} + \sigma_\eta \eta_k, & \eta_k \sim \mathcal{N}(0, s^2) \\ y_k = b \exp(\alpha_k/2) \varepsilon_k, & \varepsilon_k \sim \mathcal{N}(0, 1) \end{cases}$$

For simplicity we assume that $\alpha_1 \sim \mathcal{N}(0, 1)$.

Example: Stochastic Volatility

EM-algorithm, in this model we find four **sufficient statistics**:

$$t_1 = \sum_{t=2}^n \mathbb{E}[\alpha_t^2 | y_{1:n}]$$

$$t_2 = \sum_{t=2}^n \mathbb{E}[\alpha_t \alpha_{t-1} | y_{1:n}]$$


$$t_3 = \sum_{t=2}^n \mathbb{E}[\alpha_{t-1}^2 | y_{1:n}]$$

$$t_4 = \sum_{t=1}^n \mathbb{E}[y_t^2 \exp(-\alpha_t) | y_{1:n}]$$


The parameters are updated as,

$$a = \frac{t_2}{t_3} \quad s^2 = \frac{1}{n-1} (t_1 - \frac{t_2^2}{t_3}) \quad b^2 = \frac{t_4}{n}$$

Example: Stochastic Volatility



Example: Stochastic Volatility



Extensions to Particle Filters

Adaptive Resampling

- So far we have done **resampling every iteration**.
- This is **unnecessary**, we only need to do it if the weights have degenerated.

Adaptive Resampling

- So far we have done **resampling every iteration**.
- This is **unnecessary**, we only need to do it if the weights have degenerated.
- Use **effective sample size** (ESS) as a measure for degeneracy,

$$\text{ESS}_t = \frac{(\sum_{i=1}^N \omega_t^i)^2}{\sum_{i=1}^N (\omega_t^i)^2}$$

Adaptive Resampling

- So far we have done **resampling every iteration**.
- This is **unnecessary**, we only need to do it if the weights have degenerated.
- Use **effective sample size** (ESS) as a measure for degeneracy,

$$\text{ESS}_t = \frac{(\sum_{i=1}^N \omega_t^i)^2}{\sum_{i=1}^N (\omega_t^i)^2}$$

- If all weights are **equal** then $\text{ESS}_t = N$.

Adaptive Resampling

- So far we have done **resampling every iteration**.
- This is **unnecessary**, we only need to do it if the weights have degenerated.
- Use **effective sample size** (ESS) as a measure for degeneracy,

$$\text{ESS}_t = \frac{(\sum_{i=1}^N \omega_t^i)^2}{\sum_{i=1}^N (\omega_t^i)^2}$$

- If all weights are **equal** then $\text{ESS}_t = N$.
- If all weights except one is zero then $\text{ESS}_t = 1$.

Adaptive Resampling

- So far we have done **resampling every iteration**.
- This is **unnecessary**, we only need to do it if the weights have degenerated.
- Use **effective sample size** (ESS) as a measure for degeneracy,

$$\text{ESS}_t = \frac{(\sum_{i=1}^N \omega_t^i)^2}{\sum_{i=1}^N (\omega_t^i)^2}$$

- If all weights are **equal** then $\text{ESS}_t = N$.
- If all weights except one is zero then $\text{ESS}_t = 1$.
- Choose a **threshold** N_{ESS} , if $\text{ESS}_t < N_{\text{ESS}}$ then resample, otherwise no resampling happens.
- If no resampling happens the weights should be updated as in the SIS algorithm.

Algorithm: Particle Filter with Adaptive Resampling

Particle Filter:

Draw $\alpha_1^i \sim f(\alpha_1)$

Set $\omega_1^i = \frac{q(\alpha_1^i)g(y_1 | \alpha_1^i)}{f(\alpha_1^i)}$

Set $\Omega_1 = \sum_{i=1}^N \omega_1^i$

for $t = 2, 3, \dots, n$ **do**

 Calculate $\text{ESS}_{t-1} = \frac{(\sum_{i=1}^N \omega_t^i)^2}{\sum_{i=1}^N (\omega_t^i)^2}$

if $\text{ESS}_{t-1} < N_{\text{ESS}}$ **then**

 Draw $l^i = j$ w. pr. $\frac{\omega_{t-1}^j}{\Omega_{t-1}}$

 Set $\alpha_{1:t}^i = (\alpha_{1:t-1}^{l^i}, \alpha_t^i)$

 Set $\omega_{t-1}^i = 1/N$

end if


 Draw $\alpha_t^i \sim f(\alpha_t | \alpha_{1:t-1}^i)$

 Set $\omega_t^i = \frac{q(\alpha_t^i | \alpha_{t-1}^i)g(y_t | \alpha_t^i)}{f(\alpha_t^i | \alpha_{1:t-1}^i)} \times \omega_{t-1}^i$

 Set $\Omega_t = \sum_{i=1}^N \omega_t^i$


end for

Example: Adaptive resampling



Resampling: 43%

Example: Adaptive resampling



Resampling: 36%

Different Resampling Schemes

- We have so far only considered **multinomial resampling**.
- There are many different schemes,
 - **Residual resampling**

Different Resampling Schemes

- We have so far only considered **multinomial resampling**.
- There are many different schemes,
 - **Residual resampling**
 - **Stratified resampling**

Different Resampling Schemes

- We have so far only considered **multinomial resampling**.
- There are many different schemes,
 - **Residual resampling**
 - **Stratified resampling**
 - **Systematic resampling**

Different Resampling Schemes

- We have so far only considered **multinomial resampling**.
- There are many different schemes,
 - **Residual resampling**
 - **Stratified resampling**
 - **Systematic resampling**

Summary

- We looked at how to calculate the **likelihood** using a **particle filter**.
 - The randomness of the algorithm gives us **noisy** likelihood estimate.
 - **Solve** by performing many runs and average.

Summary

- We looked at how to calculate the **likelihood** using a **particle filter**.
 - The randomness of the algorithm gives us **noisy** likelihood estimate.
 - **Solve** by performing many runs and average.
- We looked at the **EM-algorithm**.
 - Easier if the model belongs to the **exponential family**.
 - Requires, again, the **smoothing distribution**.
 - Looked at the **fixed-lag smoother**.

Summary

- We looked at how to calculate the **likelihood** using a **particle filter**.
 - The randomness of the algorithm gives us **noisy** likelihood estimate.
 - **Solve** by performing many runs and average.
- We looked at the **EM-algorithm**.
 - Easier if the model belongs to the **exponential family**.
 - Requires, again, the **smoothing distribution**.
 - Looked at the **fixed-lag smoother**.
- We looked at some extensions regarding the **resampling**.
 - By calculating the ESS we performed **adaptive resampling** by only resampling when needed.
 - We looked at some different resampling schemes.